

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import scipy.stats as stats
```

```
1 df = pd.read_csv ("/content/student_performance_prediction.csv")
```

```
1 # configuring display options for dataframe
2 pd.set_option('display.max_columns', None) # display max columns
3 pd.set_option('display.max_rows', None) # display max rows
```

```
1 df.sample(5)
```



	Student ID	Study Hours per Week	Attendance Rate	Previous Grades	Participation in Extracurricular Activities	Parent Education Level	Passed
29460	S29461	11.8	101.8	95.9	No	Bachelor	No
189	S00190	2.4	82.0	60.4	Yes	Associate	Yes
6084	S06085	9.5	50.4	28.2	No	Bachelor	Yes
25215	S25216	13.4	82.1	62.3	No	NaN	Yes



```
1 df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 40000 entries, 0 to 39999
Data columns (total 7 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Student ID                           40000 non-null  object
 1   Study Hours per Week                 38005 non-null  float64
 2   Attendance Rate                     38008 non-null  float64
 3   Previous Grades                     38006 non-null  float64
 4   Participation in Extracurricular Activities 38000 non-null  object
 5   Parent Education Level               38000 non-null  object
 6   Passed                             38000 non-null  object
dtypes: float64(3), object(4)
memory usage: 2.1+ MB
```

```
1 print ("DataSet Shape:", df.shape) # shape
2 print ("Total Size:", df.size) # size
```



```
DataSet Shape: (40000, 7)
Total Size: 280000
```

```
1 df.columns # columns list
```



```
Index(['Student ID', 'Study Hours per Week', 'Attendance Rate',
       'Previous Grades', 'Participation in Extracurricular Activities',
       'Parent Education Level', 'Passed'],
      dtype='object')
```

✓ Handling Missing Values

```
1 # total no. of missing values
2 total_missin_vals = df.isna().sum().sum()
3 print (f"{total_missin_vals} values are missing in the dataset.")
```



```
11981 values are missing in the dataset.
```

```
1 # total % of missing values
2 missin_vals_percent = round((total_missin_vals/df.shape[0])*100)
3 print (f"{missin_vals_percent}% of values are missing in the dataset.")
```

30% of values are missing in the dataset.

```
1 cols_wid_missin_vals = df.columns[df.isna().any()].tolist()
2 print (f"List of columns with missing values: {cols_wid_missin_vals}")
```

List of columns with missing values: ['Study Hours per Week', 'Attendance Rate', 'Previous Grades', 'Participati

```
1 missin_vals_per_col = df [['Study Hours per Week', 'Attendance Rate', 'Previous Grades', 'Participation in Extracurricular Activities']]
2 missin_vals_per_col
```

	0
Study Hours per Week	1995
Attendance Rate	1992
Previous Grades	1994
Participation in Extracurricular Activities	2000
Parent Education Level	2000
Passed	2000


dtype: int64

```
1 percent_missin_vals_per_col = (missin_vals_per_col/df.shape[0])*100
2 percent_missin_vals_per_col
```

	0
Study Hours per Week	4.9875
Attendance Rate	4.9800
Previous Grades	4.9850
Participation in Extracurricular Activities	5.0000
Parent Education Level	5.0000
Passed	5.0000

dtype: float64

```
1 # visual for missing values in the data
2 import missingno as msno
3 msno.matrix(df)
4 plt.show()
```



```
1 df[['Study Hours per Week', 'Attendance Rate', 'Previous Grades']].sample(25)
```



	Study Hours per Week	Attendance Rate	Previous Grades
39166	10.6	65.3	76.5
23873	15.0	45.5	49.6
1832	14.8	95.7	83.0
38374	14.0	72.4	74.6
3806	10.2	103.4	43.0
26264	7.8	75.9	77.7
31900	3.2	45.5	54.5
15301	15.4	82.5	36.7
29081	12.3	99.4	72.1
12264	2.5	64.4	46.1
13674	8.9	94.2	54.6
15202	3.5	74.4	53.4
13680	NaN	74.7	60.5
13717	14.1	78.4	84.5
24313	17.8	81.5	52.8
27653	6.6	93.0	73.6
13989	3.4	120.4	57.7
36324	8.3	62.4	38.8
3214	13.7	80.8	69.2
877	11.1	107.6	51.8
27401	14.5	91.0	63.3
33272	16.9	81.6	NaN
19766	13.5	59.4	98.4
13918	13.9	77.0	33.4
2620	13.5	83.3	36.8



```
1 # mean interpolation for numeric columns
2
3 # Calculate the mean for each column
4 mean_values = df[['Study Hours per Week', 'Attendance Rate', 'Previous Grades']].mean()
5
6 # Fill missing values with the calculated means
7 df.fillna(mean_values, inplace=True)

1 df[['Study Hours per Week', 'Attendance Rate', 'Previous Grades']].isna().sum()
```



	0
Study Hours per Week	0
Attendance Rate	0
Previous Grades	0

dtype: int64

```
1 df[['Participation in Extracurricular Activities', 'Parent Education Level', 'Passed']].sample(25)
```



	Participation in Extracurricular Activities	Parent Education Level	Passed
23912	No	Associate	Yes
39156	Yes	Master	Yes
27364	NaN	Doctorate	No
18109	NaN	High School	No
1119	No	Bachelor	No
231	Yes	High School	No
10062	Yes	Doctorate	Yes
8388	Yes	Bachelor	Yes
4188	Yes	Associate	Yes
34485	No	Master	No
23753	No	Master	No
25530	NaN	High School	No
28458	No	High School	Yes
9388	Yes	High School	No
7863	Yes	Bachelor	No
8835	No	Bachelor	No
15635	No	Doctorate	Yes
37991	Yes	NaN	Yes
22483	No	Associate	Yes
35464	Yes	Doctorate	No
11579	Yes	Bachelor	NaN
2506	Yes	Bachelor	No
3186	No	NaN	Yes
39602	No	High School	No
30069	No	NaN	Yes



```
1 # mode interpolation for categorical columns
2
3 # 1. Select categorical columns
4 categorical_columns = df.select_dtypes(include="object").columns
5
6 # 2. Iterate over categorical columns
7 for column in categorical_columns:
8     # 3. Calculate the mode for the current column
9     mode_value = df[column].mode()[0]
10
11 # 4. Fill missing values with the mode
12 df[column] = df[column].fillna(mode_value)
```

```
1 df[['Participation in Extracurricular Activities', 'Parent Education Level', 'Passed']].isna().sum()
```



	0
Participation in Extracurricular Activities	0
Parent Education Level	0
Passed	0

dtype: int64

```
1 df.columns
```

```
Index(['Student ID', 'Study Hours per Week', 'Attendance Rate',  
      'Previous Grades', 'Participation in Extracurricular Activities',  
      'Parent Education Level', 'Passed'],  
      dtype='object')
```

✓ Handling Outliers & Skewness

✓ col 1: 'Student ID'

```
1 for col in ['Study Hours per Week', 'Attendance Rate', 'Previous Grades']:  
2     Q1 = df[col].quantile(0.25)  
3     Q3 = df[col].quantile(0.75)  
4     IQR = Q3 - Q1  
5  
6     # Define lower and upper bounds  
7     lower_bound = Q1 - 1.5 * IQR  
8     upper_bound = Q3 + 1.5 * IQR  
9  
10    # Identify outliers  
11    outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)]  
12    print(f"Outliers in {col}:")  
13    print(outliers)
```



39516	No
39557	Yes
39571	No
39576	No
39678	No
39732	Yes
39797	No
39851	Yes
39890	Yes


```

1 # Create a box plot for each numeric column
2 plt.figure(figsize=(12, 6))
3
4 plt.subplot(1, 3, 1)
5 sns.boxplot(x=df['Study Hours per Week'])
6 plt.title("Outliers in 'Study Hours per Week' Column")
7
8 plt.tight_layout()
9 plt.show()

```



Outliers in 'Study Hours per Week' Column




```

1 # Create a box plot for each numeric column
2 plt.figure(figsize=(12, 6))
3
4 plt.subplot(1, 3, 1)
5 sns.boxplot(x=df['Attendance Rate'])
6 plt.title("Outliers in 'Attendance Rate' Column")
7
8 plt.tight_layout()
9 plt.show()

```




Outliers in 'Attendance Rate' Column



```
1 # Create a box plot for each numeric column
2 plt.figure(figsize=(12, 6))
3
4 plt.subplot(1, 3, 1)
5 sns.boxplot(x=df['Previous Grades'])
6 plt.title("Outliers in 'Previous Grades' Column")
7
8 plt.tight_layout()
9 plt.show()
```




Outliers in 'Previous Grades' Column

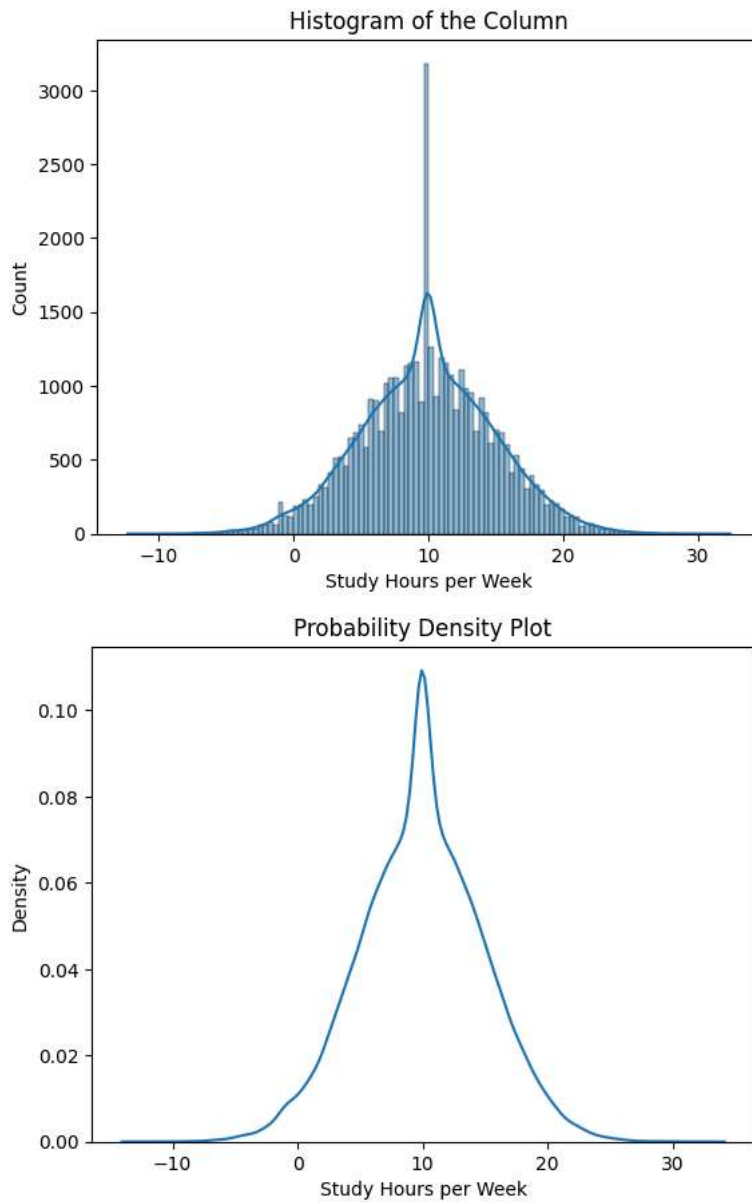


```
1 skewness = df[['Study Hours per Week', 'Attendance Rate', 'Previous Grades']].skew()
2 print("Skewness in Columns -\n", skewness)
```

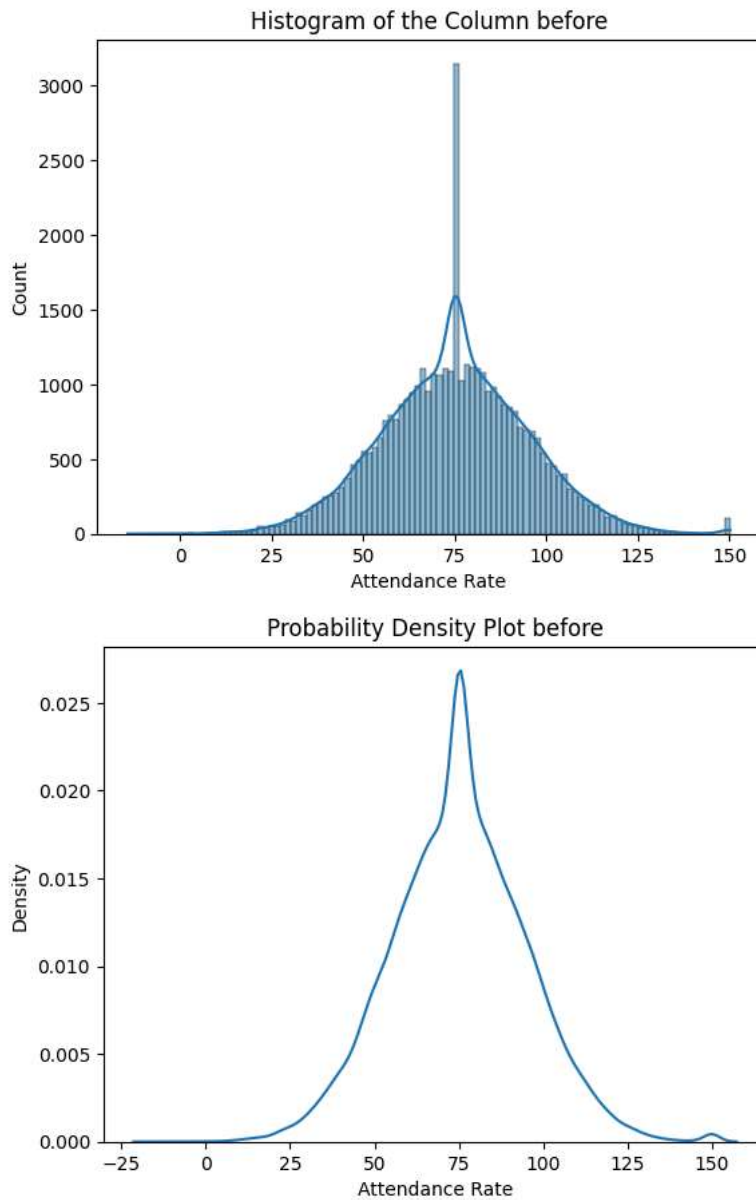


```
Skewness in Columns -
Study Hours per Week    -0.011992
Attendance Rate         0.091591
Previous Grades         1.409301
dtype: float64
```



```
1 # Histogram
2 sns.histplot(df['Study Hours per Week'], kde=True)
3 plt.title('Histogram of the Column')
4 plt.show()
5
6 # Probability Density Plot
7 sns.kdeplot(df['Study Hours per Week'])
8 plt.title('Probability Density Plot')
9 plt.show()
```



```
1 # Histogram
2 sns.histplot(df['Attendance Rate'], kde=True)
3 plt.title('Histogram of the Column before')
4 plt.show()
5
6 # Probability Density Plot
7 sns.kdeplot(df['Attendance Rate'])
8 plt.title('Probability Density Plot before')
9 plt.show()
```



```
1 # Histogram
2 sns.histplot(df['Previous Grades'], kde=True)
3 plt.title('Histogram of the Column before')
4 plt.show()
5
6 # Probability Density Plot
7 sns.kdeplot(df['Previous Grades'])
8 plt.title('Probability Density Plot before')
9 plt.show()
```



```



1 # Handling Negative Skewness and Outliers
2
3 # Handle negative values (if any):
4 df['Study Hours per Week'] = df['Study Hours per Week'].abs()
5
6 # Log transformation
7 df['Study Hours per Week'] = np.log1p(df['Study Hours per Week'])
8
9 # Define thresholds (adjust as needed)
10 lower_bound = df['Study Hours per Week'].quantile(0.05)
11 upper_bound = df['Study Hours per Week'].quantile(0.95)
12
13 # Cap outliers
14 df['Study Hours per Week'] = np.clip(df['Study Hours per Week'], lower_bound, upper_bound)

```

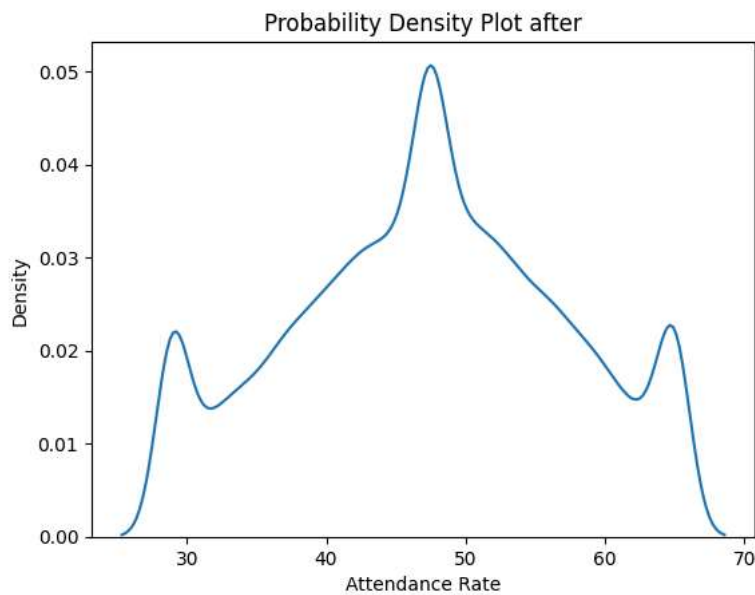

```
1 # Winsorization
2 lower_bound = df['Attendance Rate'].quantile(0.05)
3 upper_bound = df['Attendance Rate'].quantile(0.95)
4 df['Attendance Rate'] = np.clip(df['Attendance Rate'], lower_bound, upper_bound)
5
6 # Or, Box-Cox Transformation
7 from scipy.stats import boxcox
8 df['Attendance Rate'], _ = boxcox(df['Attendance Rate'] + 1) # Add 1 to handle zero values
```

```
1 # Winsorization
2 lower_bound = df['Previous Grades'].quantile(0.05)
3 upper_bound = df['Previous Grades'].quantile(0.95)
4 df['Previous Grades'] = np.clip(df['Previous Grades'], lower_bound, upper_bound)
5
6 # Or, Box-Cox Transformation
7 from scipy.stats import boxcox
8 df['Previous Grades'], _ = boxcox(df['Previous Grades'] + 1) # Add 1 to handle zero values
```


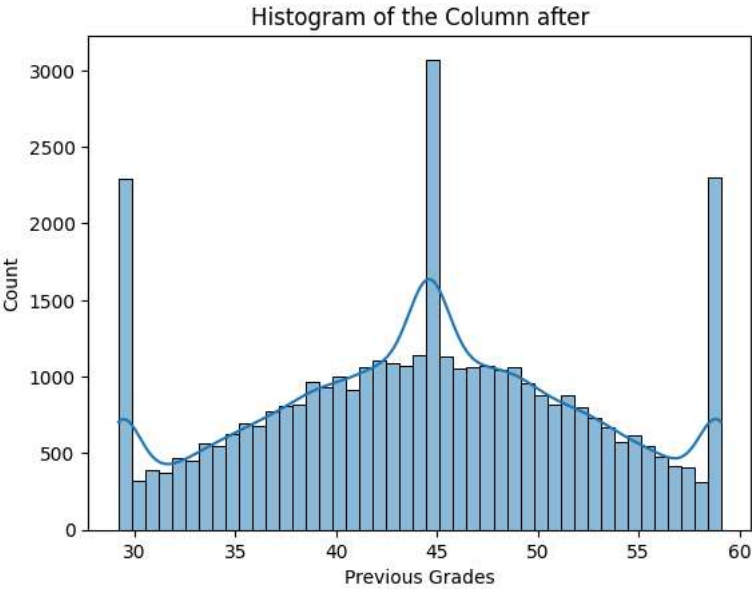
```
1 # Histogram
2 sns.histplot(df['Study Hours per Week'], kde=True)
3 plt.title('Histogram of the Column after')
4 plt.show()
5
6 # Probability Density Plot
7 sns.kdeplot(df['Study Hours per Week'])
8 plt.title('Probability Density Plot after')
9 plt.show()
```



```
1 # Histogram
2 sns.histplot(df['Attendance Rate'], kde=True)
3 plt.title('Histogram of the Column after')
4 plt.show()
5
6 # Probability Density Plot
7 sns.kdeplot(df['Attendance Rate'])
8 plt.title('Probability Density Plot after')
9 plt.show()
```



```
1 # Histogram
2 sns.histplot(df['Previous Grades'], kde=True)
3 plt.title('Histogram of the Column after')
4 plt.show()
5
6 # Probability Density Plot
7 sns.kdeplot(df['Previous Grades'])
8 plt.title('Probability Density Plot after')
9 plt.show()
```



1 df.sample(50)